



香港中文大學(深圳)
The Chinese University of Hong Kong, Shenzhen

Operating Systems

Lecture 6

System programming: sockets, pipes

Prof. Yunming XIAO
School of Data Science

Acknowledgments: Ion Stoica and Matei Zaharia, Berkeley CS 162

Recall: what's in an open file description?

- Inside kernel!
- For our purpose, the two most important things are:
 - Where to find the file data on disk
 - The current position within the file

Linux:

<https://github.com/torvalds/linux/blob/master/include/linux/fs.h>

xv6:


<https://github.com/mit-pdos/xv6-riscv/blob/riscv/kernel/file.h>

linux / include / linux / fs.h

Code Blame 3609 lines (3186 loc) · 117 KB


```
1257  */
1258  struct file {
1259      spinlock_t          f_lock;
1260      fmode_t             f_mode;
1261      const struct file_operations *f_op;
1262      struct address_space *f_mapping;
1263      void                *private_data;
1264      struct inode         *f_inode;
1265      unsigned int         f_flags;
1266      unsigned int         f_iocb_flags;
1267      const struct cred    *f_cred;
1268      struct fown_struct   *f_owner;
1269      /* --- cacheline 1 boundary (64 bytes) --- */
1270      union {
1271          const struct path f_path;
1272          struct path       __f_path;
1273      };
1274      union {
1275          /* regular files (with FMODE_ATOMIC_POS) and directories */
1276          struct mutex       f_pos_lock;
1277          /* pipes */
1278          u64                f_pipe;
1279      };
1280      loff_t               f_pos;
1281      #ifdef CONFIG_SECURITY
1282      void                *f_security;
1283      #endif
1284      /* --- cacheline 2 boundary (128 bytes) --- */
1285      errseq_t             f_wb_err;
1286      errseq_t             f_sb_err;
```

Process-specific file descriptor table inside kernel




- Suppose that we execute ``open("foo.txt")`` and that the result is 3

Process-specific file descriptor table inside kernel




- Suppose that we execute ``open("foo.txt")`` and that the result is 3
- Next, suppose that we execute ``read(3, buf, 100)`` and the result is 100

Process-specific file descriptor table inside kernel




- Suppose that we execute ``open("foo.txt")`` and that the result is 3
- Next, suppose that we execute ``read(3, buf, 100)`` and the result is 100

Process-specific file descriptor table inside kernel




- Suppose that we execute ``open("foo.txt")`` and that the result is 3
- Next, suppose that we execute ``read(3, buf, 100)`` and the result is 100
- Finally, suppose that we execute ``close(3)``

Process-specific file descriptor table inside kernel




- Suppose that we execute ``open("foo.txt")`` and that the result is 3
- Next, suppose that we execute ``read(3, buf, 100)`` and the result is 100
- Finally, suppose that we execute ``close(3)``

Instead of closing, let's fork()!




Open file description is *aliased*

`read(3, buf, 100)`



Open file description is *aliased*

`read(3, buf, 100)`




Open file description is *aliased*

`read(3, buf, 100)`

Process 1

`read(3, buf, 100)`

Process 2



Not shown: Initially contains 0, 1, and 2 (stdin, stdout, stderr)


Open file description is *aliased*

`read(3, buf, 100)`

Process 1

`read(3, buf, 100)`

Process 2



Not shown: Initially contains 0, 1, and 2 (stdin, stdout, stderr)


File description is *copied*

`read(3, buf, 100)`
`close(3)`

Process 1

`read(3, buf, 100)`

Process 2



Not shown: Initially contains 0, 1, and 2 (stdin, stdout, stderr)


File description is copied

`read(3, buf, 100)`
`close(3)`

Process 1

`read(3, buf, 100)`

Process 2



Not shown: Initially contains 0, 1, and 2 (stdin, stdout, stderr)

- Open file description remains alive until no file descriptors in any process refer to it

End of recall


Why is aliasing the open file description a good idea?

- It allows for *shared resources* between processes

Example: shared terminal emulator


- When you `fork()` a process, the parent's and child's `printf` outputs go to the same terminal

Example: shared terminal emulator




Example: shared terminal emulator

`close(0)`



Example: shared terminal emulator

`close(0)`



- If one process closes stdin (0), it remains open in other processes

Inter-process communication (IPC)

Communication between processes

- Can we view files as communication channels?

```
write(wfd, wbuf, wlen);
```




```
n = read(rfd, rbuf, rmax);
```

- Producer and Consumer of a file may be distinct processes
 - May be separated in time (or not)
- However, what if data written once and consumed once?
 - Don't we want something more like a queue?
 - Can still look like file I/O!

Communication across the world looks like file I/O too!

```
write(wfd, wbuf, wlen);
```



```
n = read(rfd, rbuf, rmax);
```


- Connected queues over the Internet
 - But what's the analog of open?
 - What is the namespace?
 - How are they connected in time?

Request response protocol

Client (issues requests)

Server (performs operations)

```
write(rqfd, rqbuf, buflen);
```



```
n = read(rfd, rbuf, rmax);
```

```
write(wfd, respbuf, len);
```


```
n = read(resfd, resbuf, resmax);
```

Request response protocol: across network

Client (issues requests)

Server (performs operations)


```
write(rqfd, rqbuf, buflen);
```



The socket abstraction: endpoint for communication

- Key idea: communication across the world looks like file I/O

```
write(wfd, wbuf, wlen);
```



```
n = read(rfd, rbuf, rmax);
```

- Sockets: endpoint for communication
 - Queues to temporarily hold results
- Connection: two sockets connected over the network → IPC over network!
 - How to open()?
 - What is the namespace?
 - How are they connected in time?

Sockets: more details

- Socket: an abstraction for one endpoint of a network connection
 - Another mechanism for inter-process communication
 - Most operating systems (Linux, Mac OS X, Windows) provide this, even if they don't copy rest of UNIX I/O
 - Standardized by POSIX
- First introduced in 4.2 BSD (Berkeley Standard Distribution) Unix
- Same abstraction for any kind of network
 - Local (within same machine)
 - The Internet (TCP/IP, UDP/IP)
 - Things “no one” uses anymore (OSI, Appletalk, IPX, ...)

Sockets: more details

- Looks just like a file with a file descriptor
 - Corresponds to a network connection (two queues)
 - write adds to output queue (queue of data destined for other side)
 - read removes from it input queue (queue of data destined for this side)
 - Some operations do not work, e.g. lseek
- How can we use sockets to support real applications?
 - A bidirectional byte stream isn't useful on its own...
 - May need messaging facility to partition stream into chunks
 - May need RPC facility to translate one environment to another and provide the abstraction of a function call over the network

Simple example: echo server



Simple example: echo server

Client (issues requests)

Server (serves requests)


```
fgets(sndbuf, bufsize, stdin);
```

```
write(sockfd, sndbuf, strlen(sndbuf)+1);
```

```
n = read(sockfd, reqbuf, ...);
```


```
n = read(sockfd, rcvbuf, ...);
```

```
write(sockfd, reqbuf, ...);
```



Echo client-server example

```
void client(int sockfd) {
    int n;
    char sndbuf[MAXIN]; char rcvbuf[MAXOUT];
    while (1) {
        fgets(sndbuf, MAXIN, stdin);          /* prompt */
        write(sockfd, sndbuf, strlen(sndbuf)+1); /* send (including null terminator) */
        memset(rcvbuf, 0, MAXOUT);            /* clear */
        n=read(sockfd, rcvbuf, MAXOUT);        /* receive */
        write(STDOUT_FILENO, rcvbuf, n);       /* echo */
    }
}
```



```
void server(int consockfd) {
    char reqbuf[MAXREQ];
    int n;
    while (1) {
        memset(reqbuf, 0, MAXREQ);
        n = read(consockfd, reqbuf, MAXREQ); /* recv */
        if (n <= 0) return;
        write(STDOUT_FILENO, reqbuf, n);
        write(consockfd, reqbuf, n); /* echo */
    }
}
```

What assumptions are we making?

- Reliable
 - Write to a file => Read it back. Nothing is lost.
 - Write to a (TCP) socket => Read from the other side, same.
- In order (sequential stream)
 - Write X then write Y => read gets X then read gets Y
- When ready?
 - File read gets whatever is there at the time
 - Actually, need to loop and read until we receive the terminator ('\0')
 - Assumes writing already took place
 - Blocks if nothing has arrived yet

Socket creation

- File systems provide a collection of permanent objects in a structured name space:
 - Processes open, read/write/close them
 - Files exist independently of processes
 - Easy to name what file to open()
- Pipes: one-way communication between processes on same (physical) machine (see later)
 - Single queue
 - Created transiently by a call to pipe()
 - Passed from parent to children (descriptors inherited from parent process)


Socket creation

- Sockets: two-way communication between processes on same or different machine
 - Two queues (one in each direction)
 - Processes can be on separate machines: no common ancestor
 - How do we name the objects we are **opening**?
 - How do these completely independent programs know that the other wants to “talk” to them?

Namespaces for communication over IP


- Hostname
 - yunmingxiao.github.io
- IP address
 - 185.199.111.153 (IPv4, 32-bit Integer)
 - 2606:50c0:8000::153 (IPv6, 128-bit Integer)
- Port Number
 - 0-1023 are “[well known](#)” or “system” ports
 - E.g., 80 (http), 443 (secure web)
 - Superuser privileges to bind to one
 - 1024 – 49151 are “registered” ports ([registry](#))
 - Assigned by IANA for specific services
 - 49152–65535 ($2^{15}+2^{14}$ to $2^{16}-1$) are “dynamic” or “private”
 - Automatically allocated as “ephemeral ports”

Connection setup over TCP/IP




- Special kind of socket: server socket
 - Has file descriptor
 - Can't read or write
- Two operations:
 - `listen()`: Start allowing clients to connect
 - `accept()`: Create a new socket for a particular client

Connection setup over TCP/IP




- 5-tuple identifies each connection:
 1. Source IP address
 2. Destination IP address
 3. Source port number
 4. Destination port number
 5. Protocols (always TCP here)
- Often, client port “randomly” assigned
 - Done by OS during client socket setup
- Server port often “well known”
 - 80 (web), 443 (secure web), 25 (sendmail), etc
 - Well-known ports from 0 to 1023

Web server




Client-server models



- File servers, web, FTP, Databases, ...
- Many clients accessing a common server

Simple web server



Client code

```
char *host_name, *port_name;

// Create a socket
struct addrinfo *server = lookup_host(host_name, port_name);
int sock_fd = socket(server->ai_family, server->ai_socktype,
                     server->ai_protocol);

// Connect to specified host and port
connect(sock_fd, server->ai_addr, server->ai_addrlen);

// Carry out client-server protocol
run_client(sock_fd);

// Clean up on termination
close(sock_fd);
```

Client-side: getting the server address

```
struct addrinfo *lookup_host(char *host_name, char *port_name) {
    struct addrinfo *server;
    struct addrinfo hints;
    memset(&hints, 0, sizeof(hints));
    hints.ai_family = AF_UNSPEC;          /* Includes AF_INET and AF_INET6 */
    hints.ai_socktype = SOCK_STREAM;     /* Essentially TCP/IP */

    int rv = getaddrinfo(host_name, port_name, &hints, &server);
    if (rv != 0) {
        printf("getaddrinfo failed: %s\n", gai_strerror(rv));
        return NULL;
    }
    return server;
}
```

Server code (v1)

```
// Create socket to listen for client connections
char *port_name;
struct addrinfo *server = setup_address(port_name);
int server_socket = socket(server->ai_family, server->ai_socktype,
                           server->ai_protocol);

// Bind socket to specific port
bind(server_socket, server->ai_addr, server->ai_addrlen);
// Start listening for new client connections
listen(server_socket, MAX_QUEUE);

while (1) {
    // Accept a new client connection, obtaining a new socket
    int conn_socket = accept(server_socket, NULL, NULL);
    serve_client(conn_socket);
    close(conn_socket);
}
close(server_socket);
```

Server address: itself (wildcard IP), passive


```
struct addrinfo *setup_address(char *port) {
    struct addrinfo *server;
    struct addrinfo hints;
    memset(&hints, 0, sizeof(hints));
    hints.ai_family = AF_UNSPEC; /* Includes AF_INET and AF_INET6 */
    hints.ai_socktype = SOCK_STREAM; /* Essentially TCP/IP */
    hints.ai_flags = AI_PASSIVE; /* Set up for server socket */
    int rv = getaddrinfo(NULL, port, &hints, &server); /* No address! */
    if (rv != 0) {
        printf("getaddrinfo failed: %s\n", gai_strerror(rv));
        return NULL;
    }
    return server;
}
```

- Accepts any connections on the specified port

How could the server protect itself?

- Handle each connection in a separate process
 - This will mean that the logic serving each request will be “sandboxed” away from the main server process
- In the following code, keep in mind:
 - `fork()` will duplicate all of the parent’s file descriptors (i.e. pointers to sockets!)
 - We keep control over accepting new connections in the parent
 - New child connection for each remote client

Server with protection (each process has own process)




Server code (v2)

```
// Socket setup code elided...
listen(server_socket, MAX_QUEUE);
while (1) {
    // Accept a new client connection, obtaining a new socket
    int conn_socket = accept(server_socket, NULL, NULL);
    pid_t pid = fork();
    if (pid == 0) {
        close(server_socket);
        serve_client(conn_socket);
        close(conn_socket);
        exit(0);
    } else {
        close(conn_socket);
        wait(NULL);
    }
}
close(server_socket);
```

How to make a concurrent server

- So far, in the server:
 - Listen will queue requests
 - Buffering present elsewhere
 - But server *waits* for each connection to terminate before serving the next
 - This is the standard shell pattern
- A concurrent server can handle and service a new connection before the previous client disconnects
 - Simple – just don't wait in parent!
 - Perhaps not so simple – multiple child processes better not have data races with one another through file system/etc!

Server with protection (each process has own process)




Server code (v3)

```
// Socket setup code elided...
listen(server_socket, MAX_QUEUE);
while (1) {
    // Accept a new client connection, obtaining a new socket
    int conn_socket = accept(server_socket, NULL, NULL);
    pid_t pid = fork();
    if (pid == 0) {
        close(server_socket);
        serve_client(conn_socket);
        close(conn_socket);
        exit(0);
    } else {
        close(conn_socket);
        // wait(NULL);
    }
}
close(server_socket);
```

Faster concurrent server (without protection)


- Spawn a new *thread* to handle each connection
 - Lower overhead spawning process (less to do)
- Main *thread* initiates new client connections without waiting for previously spawned threads
- Why give up the protection of separate processes?
 - More efficient to create new threads
 - More efficient to switch between threads
- Even more potential for data races (need synchronization?)
 - Through shared memory structures
 - Through file system

Server with concurrency, without protection



Thread pools: more later!

- Problem with previous version: unbounded threads
 - When web-site becomes too popular – throughput sinks
 - Instead, allocate a bounded “pool” of worker threads, representing the maximum level of multiprogramming



```
master() {  
    allocThreads(worker, queue);  
    while (1) {  
        conn = AcceptCon();  
        Enqueue(queue, conn);  
        wakeup(queue);  
    }  
}
```

```
worker(queue) {  
    while (1) {  
        conn = Dequeue(queue);  
        if (conn == NULL)  
            sleepOn(queue);  
        else  
            ServeWebPage(conn);  
    }  
}
```

Now let's talk about pipe

- A different IPC mechanism other than socket

Single-process pipe example (not that interesting yet...)

```
#include <unistd.h>

int main(int argc, char *argv[]) {
    char *msg = "Message in a pipe.\n";
    char buf[BUFSIZE];

    int pipe_fd[2];
    if (pipe(pipe_fd) == -1) {
        fprintf(stderr, "Pipe failed.\n");
        return EXIT_FAILURE;
    }
    ssize_t writelen = write(pipe_fd[1], msg, strlen(msg)+1);
    printf("Sent: %s [%ld, %ld]\n", msg, strlen(msg)+1, writelen);


    ssize_t readlen = read(pipe_fd[0], buf, BUFSIZE);
    printf("Rcvd: %s [%ld]\n", msg, readlen);

    close(pipe_fd[0]);
    close(pipe_fd[1]);
}
```


Could be useful for
multithreaded processes...

- Write to pipe_fd[1]
- Read from pipe_fd[0]

Example: pipes between processes



Example: channel from parent to child



Inter-process communication (IPC): parent → child

```
// continuing from earlier
pid_t pid = fork();
if (pid < 0) {
    fprintf (stderr, "Fork failed.\n");
    return EXIT_FAILURE;
}
if (pid != 0) {
    close(pipe_fd[0]); // Not using this descriptor!
    ssize_t writelen = write(pipe_fd[1], msg, msglen);
    printf("Parent: %s [%ld, %ld]\n", msg, msglen, writelen);
} else {
    close(pipe_fd[1]); // Not using this descriptor!
    ssize_t readlen = read(pipe_fd[0], buf, BUFSIZE);
    printf("Child Rcvd: %s [%ld]\n", msg, readlen);
}
```

Discussion

- Why pipe over socket?
 - Socket is more general and heavier
 - Connection setup
 - Addressing (IP, port, etc)
 - Protocol machinery (e.g., TCP congestion control)
 - Socket is not fork-native
 - Parent/child must explicitly connect
- Why pipe over local files?
 - Files require a pathname (global namespace)
 - Files are persistent, pipes are ephemeral
 - Files require explicit synchronization
 - File offsets, locking, races
- Essentially, pipe is **fork-native**

Takeaway

- UNIX provides multiple communication mechanisms

Scope	Mechanism
Same thread	Function calls
Across threads	Shared memory
Parent-child processes	Pipe
Same machine, unrelated processes	UNIX domain socket / files (named shared memory)
Across machines	Socket

Conclusion

- Recall: Everything is a file!
 - `open()`, `read()`, `write()`, and `close()` used for wide variety of I/O:
 - Devices (terminals, printers, etc.)
 - Regular files on disk
 - Networking (sockets)
 - Local inter-process communication (pipes, sockets)
- Processes have two parts
 - Threads (Concurrency)
 - Address Spaces (Protection)
- Various textbooks talk about *processes*
 - When this concerns concurrency, really talking about thread portion of a process
 - When this concerns protection, talking about address space portion of a process
- Stack is essential part of computation
 - Every thread has two stacks: user-level (in address space) and kernel
 - The kernel stack + support often called the “kernel thread”

Please help to make this class better!



<https://forms.office.com/r/UcqDW4fpdq>

Thanks